



Image classification of fashion items

Agenda

- Introduction
- Methodology
- Results
- Demo
- Conclusions



Introduction

- Fashion-MNIST is a data set of 28x28 grayscale images representing fashion items.
- Due to its complexity, it serves as a benchmark for evaluating and comparing image classification algorithms.
- Fashion item classification has practical applications in industries such as e-commerce, retail, and fashion.
- The data set contains 60,000 training images, 10,000 test images across 10 fashion categories.





Methodology

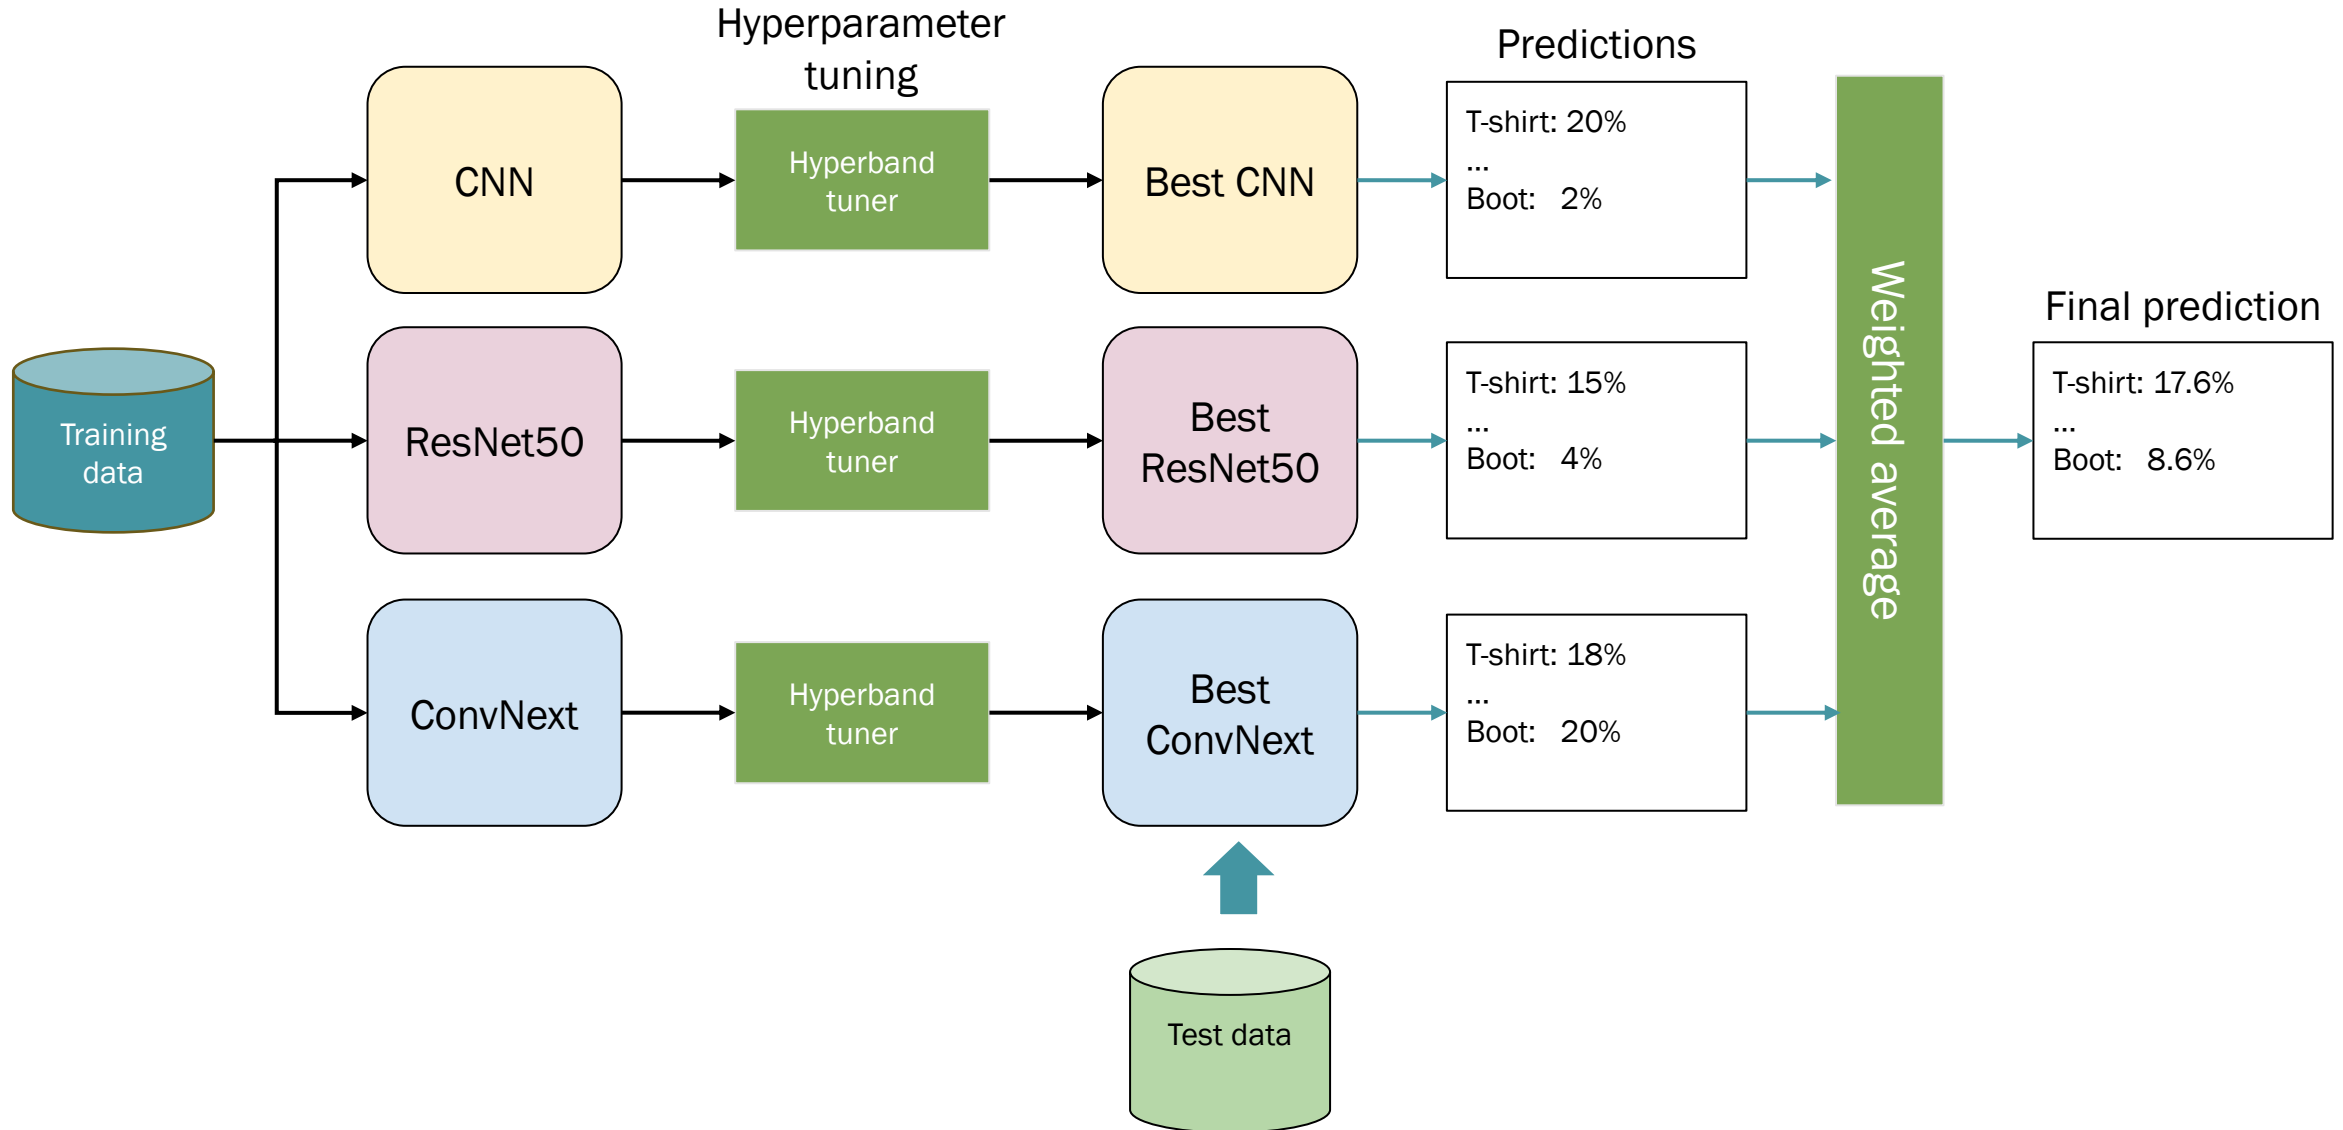
Architecture:

- Model ensemble of various neural network architectures.

Training:

- Hyperparameter tuning.
- Data augmentation.

Overview



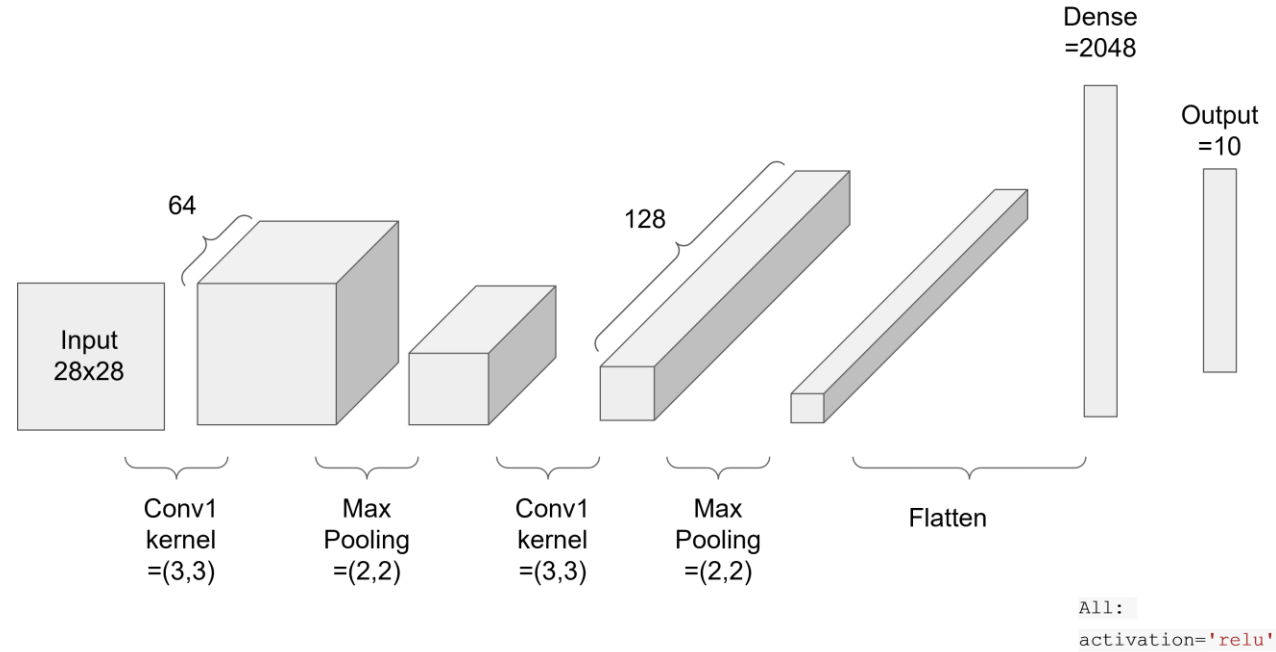
Convolutional Neural Network

CNNs are the most popular deep learning models used to extract features from images.

In our case, a custom CNN was designed. The final architecture consists of roughly 5 layers shown on the right.

The hyperparameters that need to be tuned include:

- optimizer, learning rate, beta 1 & 2, weight decay, and dropout rate.



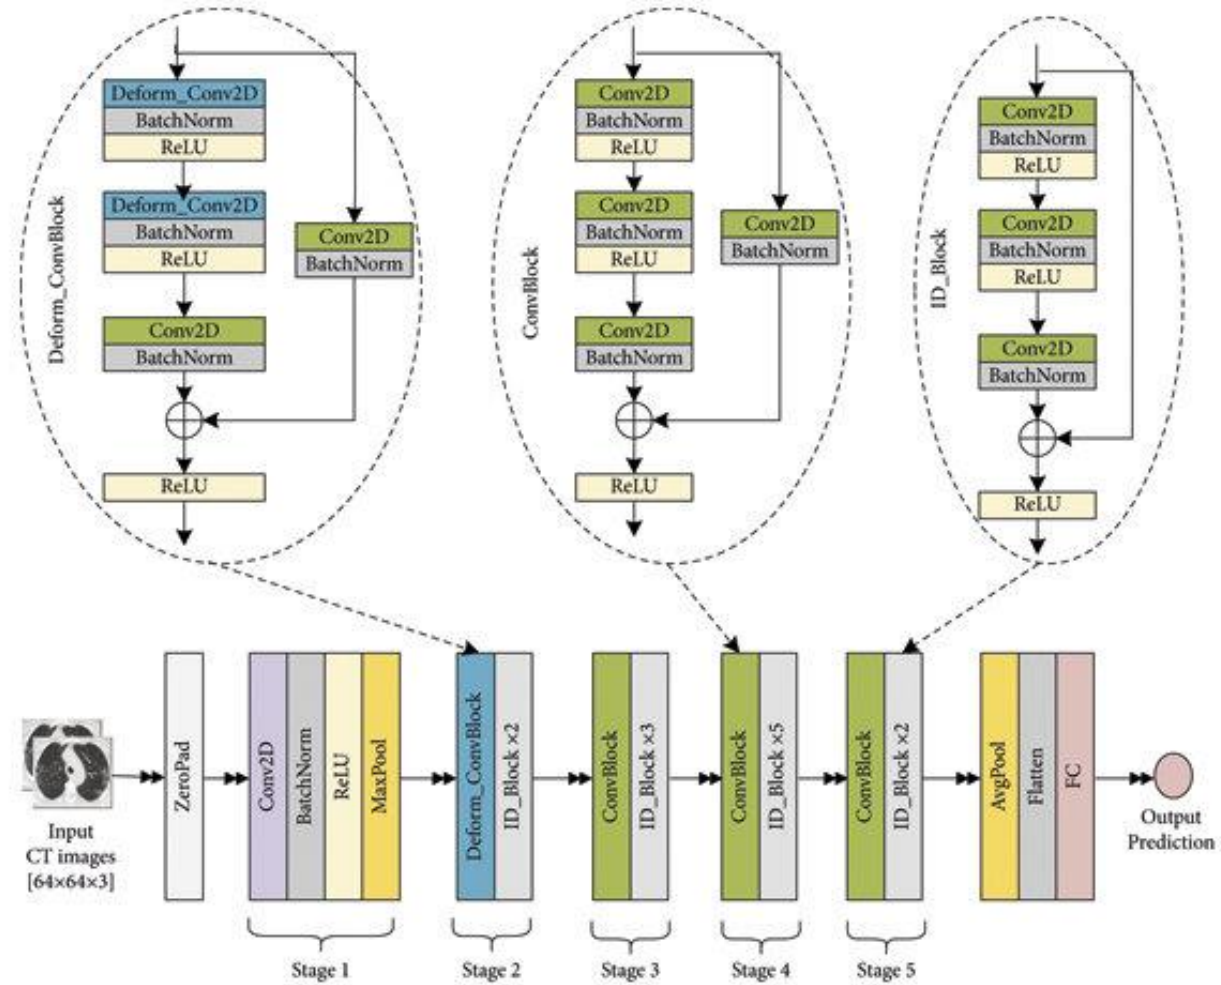
ResNet50

ResNet50 is a specific CNN architecture known for its residual connections between different stages.

In this work, we trained a ResNet50 from scratch to adapt the architecture to the problem at hand.

The hyperparameters that need to be tuned include:

- optimizer, learning rate, beta 1 & 2, weight decay, and dropout rate.



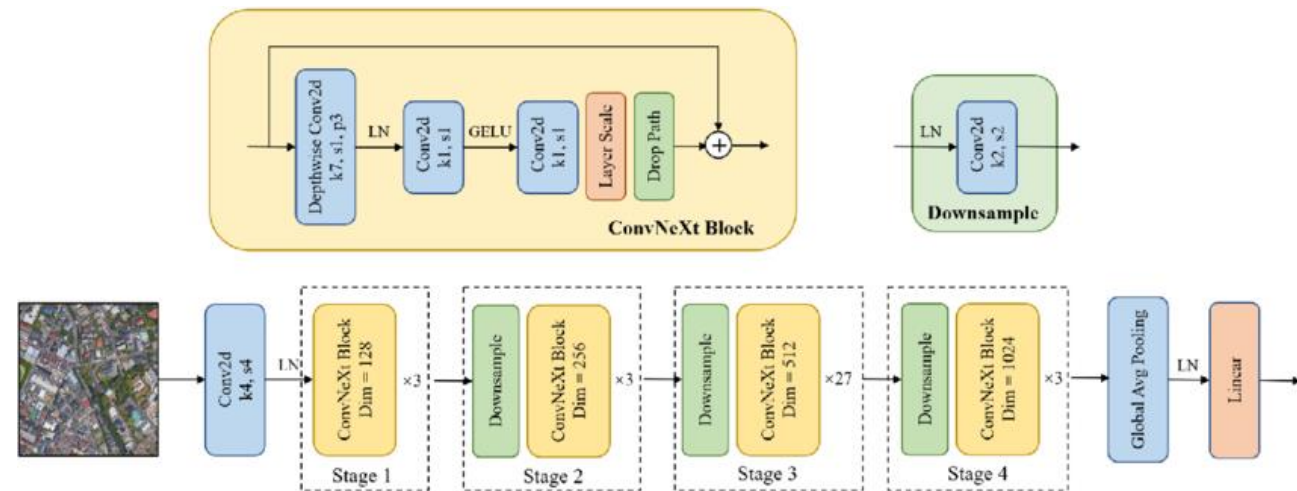
ConvNeXt

ConvNeXt is a ‘modernization’ of a standard ResNet toward the design of a vision Transformer (ViT).

Transfer learning was applied to this model, for comparison with the other models which were trained from scratch.

The hyperparameters that need to be tuned include:

- optimizer, learning rate, fine tuning layer, weight decay, and dropout rate.



Hyperparameter tuning using Hyperband


Basics:

- Hyperparameter tuning posed as a pure-exploration multi-armed bandit problem.
- Improves on the idea of successive halving.

Characteristics:

- Allocates resources (epochs) to the most promising configurations, discards the rest

Successive halving



Could we be discarding good candidates too early?

Hyperparameter tuning using Hyperband

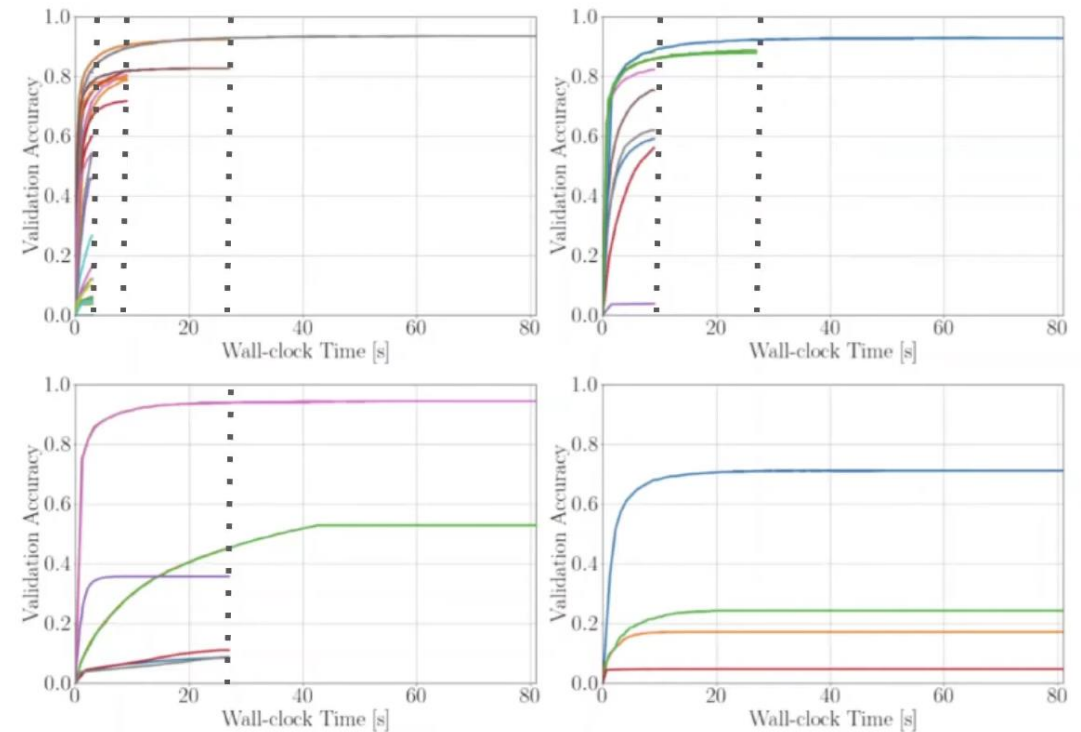
Basics:

- Hyperparameter tuning posed as a pure-exploration multi-armed bandit problem.
- Improves on the idea of successive halving.

Characteristics:

- Allocates resources (epochs) to the most promising configurations, discards the rest
- Considers that good candidates might be discarded early so it uses different brackets.

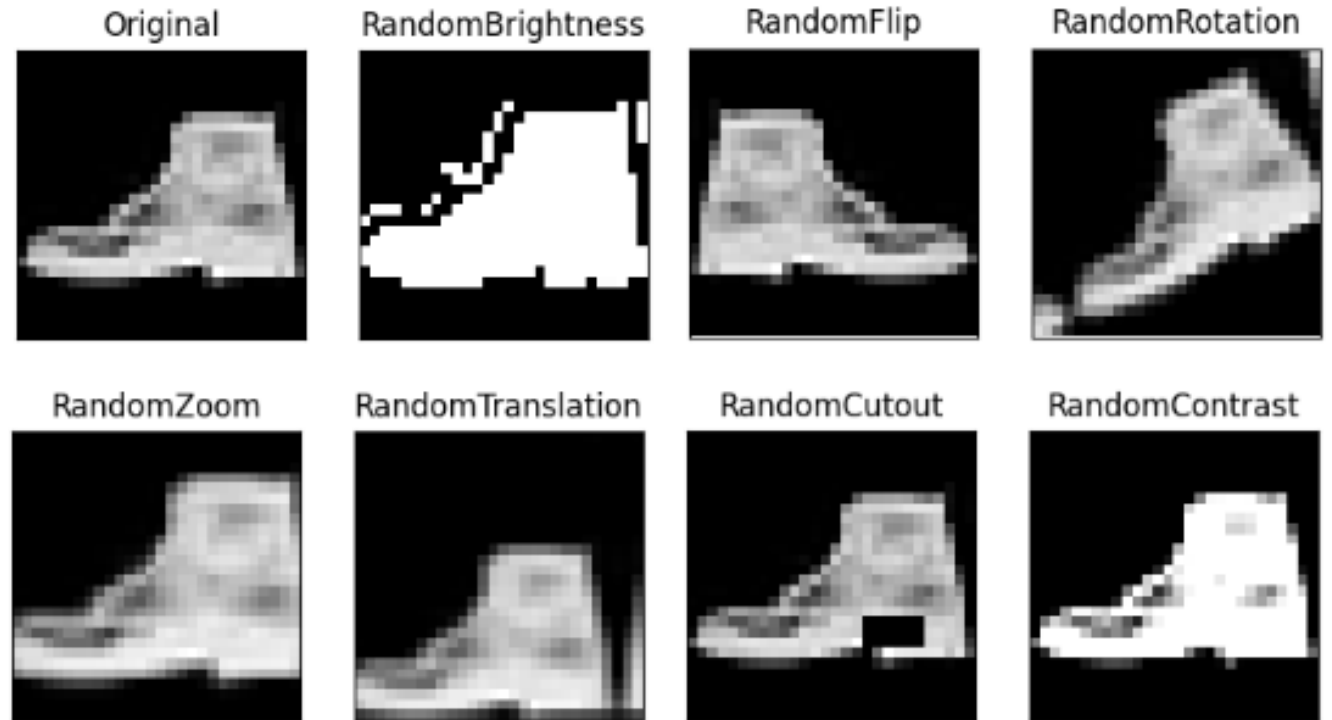
Hyperband brackets



Improving robustness through data augmentation

The following transformations to the input data were considered:

- Random Brightness
- Random Flip
- Random Rotation
- Random Zoom
- Random Translation
- Random Cutout
- Random Contrast
- Synthetic image generation (GAN)*




Improving training with learning rate reduction on Plateau

Adapting learning rate through training can help improve the model's performance and avoid local minima.

A common approach is to reduce the learning rate once the performance stops improving by a few epochs (called patience).

We use this technique with a patience of 1 epoch.





Results



Simple CNN performance

Summary

Training accuracy	Validation accuracy	Testing accuracy
0.9358	0.9107	0.9016


Best parameters

```
'optimizer': 'nadam',  
'learning_rate': 0.001,  
'beta_1': 0.9,  
'beta_2': 0.99,  
'weight_decay': 0.0,  
'dropout': 0.1
```



Learning curves

Performance on testing set





ResNet50 performance

Summary

Training accuracy	Validation accuracy	Testing accuracy
0.9622	0.9104	0.9006

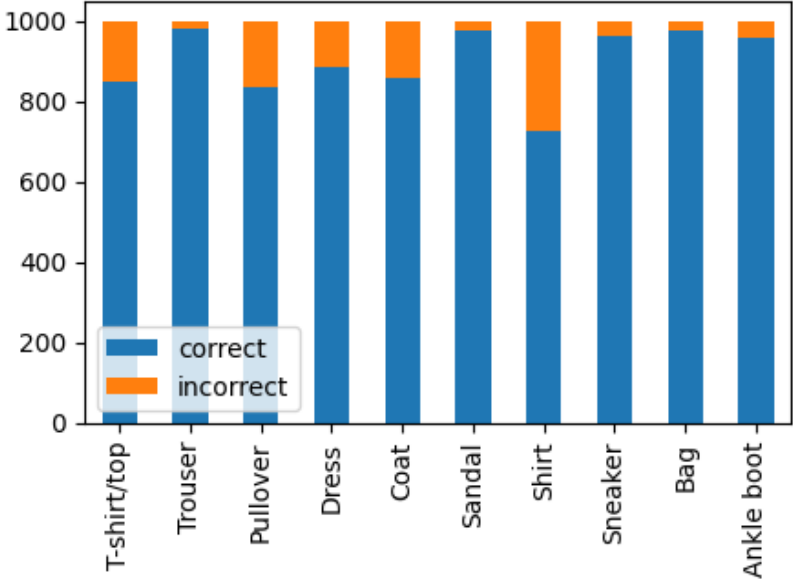
Best parameters

```
'optimizer': 'nadam',  
'learning_rate': 0.01,  
'beta_1': 0.9,  
'beta_2': 0.9999,  
'weight_decay': 1e-05,  
'dropout': 0.1,
```



Learning curves

Performance on testing set



ConvNeXt performance


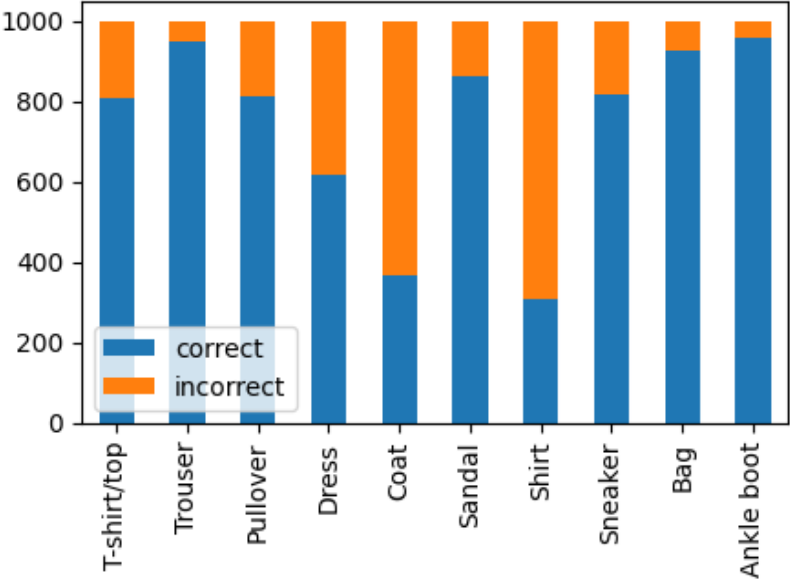
Summary

Training accuracy	Validation accuracy	Testing accuracy
0.7459	0.7425	0.7423

Best parameters

```
'optimizer': 'nadam',  
'learning_rate': 0.001,  
'fine_tune_at': 143,  
'weight_decay': 1e-04,  
'lr_schedule':  
94_step_0_9  
'dropout': 0.1,
```


Performance on testing set



Learning curves

Ensemble performance on testing set

Model	Testing set	Weight
CNN	0.9016	0.45
ResNet50	0.9007	0.45
ConvNeXt	0.7423	0.10
Ensemble	0.911	




Our model confuses T-shirts with Shirts the most



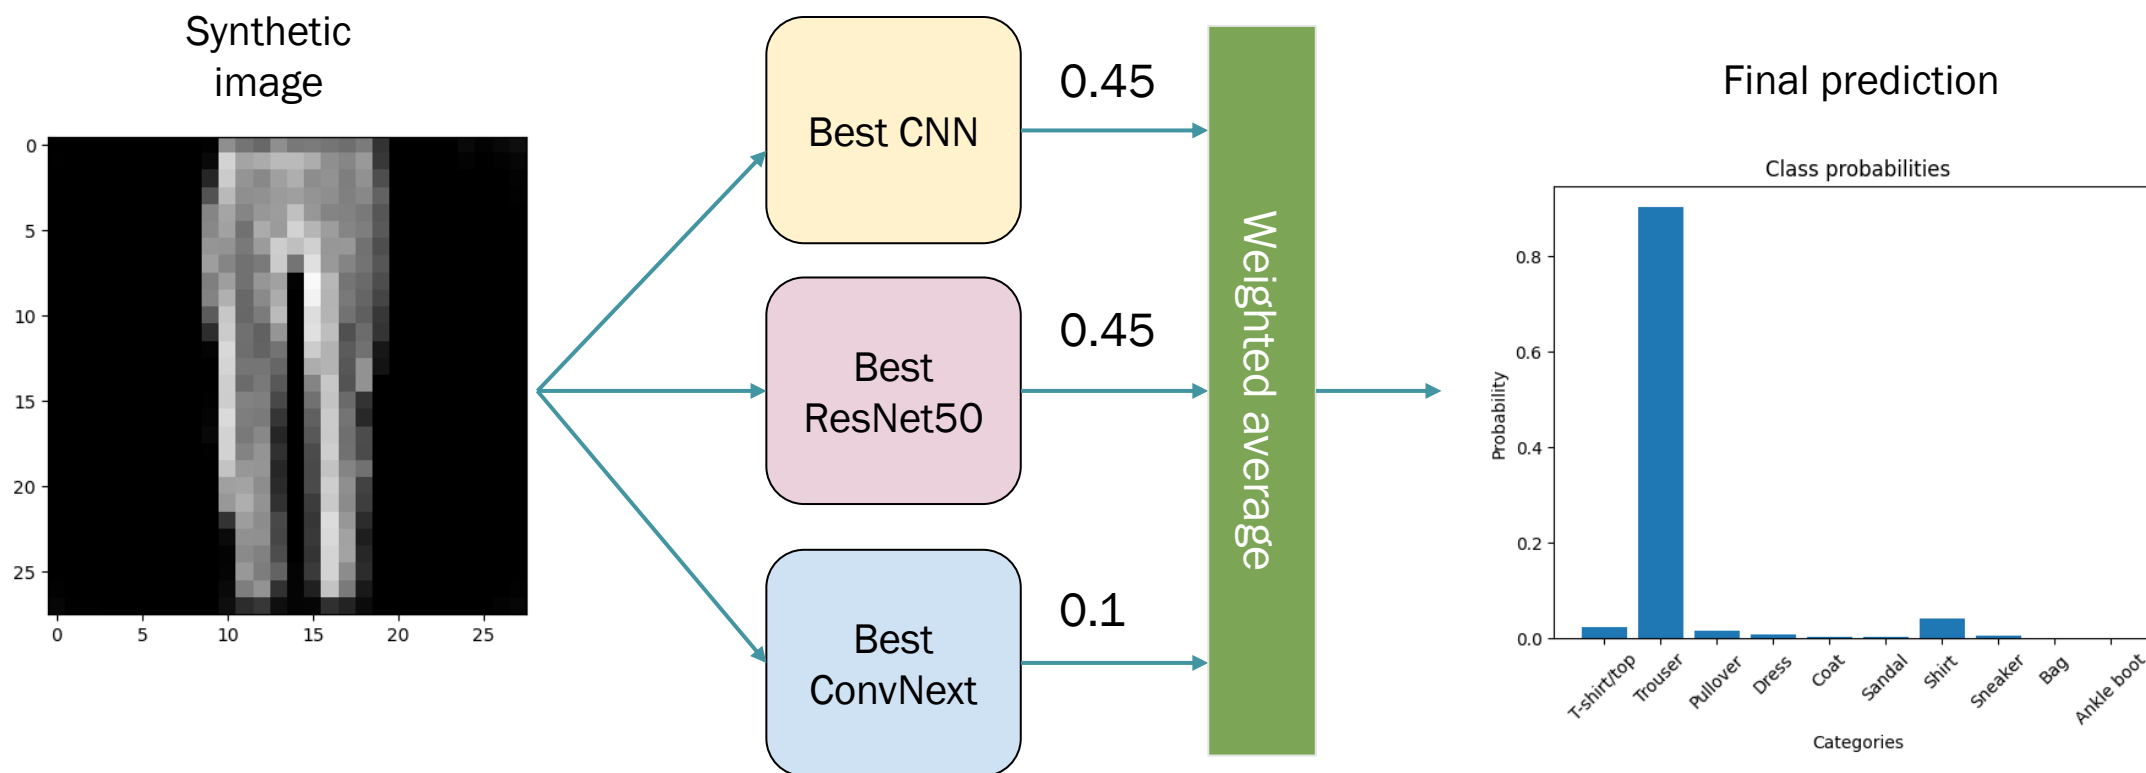
DEMO

Performance against unseen fashion item
samples

Synthetic image generation using GAN



Demo



Conclusions

- Hyperparameter tuning through Hyperband is useful for balancing resources and training time.
- Training models from scratch appears to be better than using transfer learning likely due to the differences in the images used for pretraining models.
- Data augmentation through GAN needs to be further tuned to create 'difficult' samples otherwise the benefits are small.
- The model confuses t-shirts with shirts the most, for some samples it is even hard for a human to distinguish.
- Model ensemble provides slightly better performance but for this problem a traditional CNN could be a simpler and viable approach.

Questions?